

Demo Problem 1: Hierarchical Clustering

The data set `polls.txt` contains voting data of 11 different regions from the 2017 municipal elections. The variables are the percentages of votes each of Finland's eight largest political parties gained in the elections. We will use clustering methods to see which regions resemble each other most closely.

- a) Scatter plot the variables. Can you spot the different clusters?
- b) Calculate the euclidean distances between the countries.
- c) Perform the “bottom up” hierarchical clustering by hand. Aggregate two clusters using the minimum distance (single linkage).
- d) Repeat (c) using the function `hclust`.
- e) Plot the classification tree (dendrogram).
- f) Repeat the steps by aggregating the clusters using the average (average linkage) and the maximum (complete linkage). Compare the results.
- g) Where would you cut the tree?

Demo Problem 2: k -Means Clustering

Use the data `bank.txt`. The first column contains the true classification.

- a) Apply the k -means algorithm to obtain 2 clusters.
- b) How many observations are classified to a wrong category?
- c) Change the seed number and see if it affects the results.
- d) Apply the k -means algorithm to obtain 3 clusters. Does the seed number affect the results here?

Homework Problem 1: Hierarchical Clustering

Repeat steps (a)-(b) and (d)-(g) of Problem 1 for the iris data set. The data set can be accessed from the package `MASS` via the command: `data(iris)`. Leave out the variable `species`. Remember to provide figures related to (a), (e) and (f). In addition, answer the following:

- h) Does the seed number affect the results in hierarchical clustering? Justify.
- i) Which aggregation metric (single, complete, average) results in the best separation of the different species? Justify your answer using the dendrogram.